

# ABDULLAH KASSAS

Data & AI Architect

+1(437)-988-3913 [a.kassas@gmail.com](mailto:a.kassas@gmail.com) [linkedin.com/in/3bdallah](https://linkedin.com/in/3bdallah) [github.com/i3bdallah](https://github.com/i3bdallah) Toronto, Canada  
English, Arabic Canadian

## Professional Summary

Data & AI Architect with 13+ years of experience designing and delivering end-to-end data and AI solutions on cloud platforms (Azure, AWS, Databricks, Fabric). Specializing in lakehouse architectures, scalable data platforms, and enterprise data governance. Leveraging Generative AI tools and agentic coding workflows to accelerate data platform development and engineering delivery.

## Experience

### Jarvis Consulting Group

Dec 2025 – Present

Toronto, Canada

Manager, Data Engineering

- Advising major Canadian banks on Databricks lakehouse adoption, designing scalable data platform architectures and leading legacy-to-lakehouse migration strategies for enterprise financial data workloads within highly regulated environments.
- Leading technical delivery of Databricks implementations including Unity Catalog modernization for centralized governance, access control, data lineage, and medallion architecture design across banking clients.
- Providing strategic guidance on cloud data modernization and GenAI integration, ensuring alignment with regulatory compliance, data privacy requirements, and financial services security standards.

### Gore Mutual Insurance

Feb 2024 – Nov 2025

Toronto, Canada

Lead Data Engineer

- Architected a modern data platform on Databricks lakehouse medallion architecture with Azure services (ADF, SQL Server, ADLS, Event Hubs), enabling 10+ InsurTech data products and AI models.
- Built 30+ PySpark data pipelines for ingestion, transformation, and modeling across the medallion layers, enforcing data quality checks and Delta Lake best practices.
- Designed and developed a Master Data Management (MDM) solution to govern critical insurance data entities, ensuring consistency across 5 source systems integrated into the lakehouse.
- Leveraged GenAI and ML within Databricks Unity Catalog to automate data classification and PII tagging, reducing manual governance effort and improving sensitive data handling.
- Automated the data platform's CI/CD pipelines for build, release, testing, and production deployment, reducing time-to-insight by 30% and minimizing human error.
- Led agile delivery for the data team, mentoring 4 junior data engineers while managing sprints, retrospectives, and backlog grooming to ensure timely releases.

### Avenue One

Dec 2022 – Nov 2023

New York, USA (Remote)

Senior Data Engineer

- Established a process mining framework using Celonis, driving measurable operational efficiency improvements across 3 key real-estate workflows.
- Partnered with Six Sigma experts to optimize real estate processes, leveraging data analytics to identify bottlenecks and reduce cycle times.
- Developed interactive BI dashboards enabling real-time KPI monitoring and alerting for the investment platform, covering 7 key metrics.
- Spearheaded a data-driven diagnosis of home renovation cost inefficiencies, improving house valuation model accuracy by 25% and delivering over \$300K in savings.

### BizzyCar

Feb 2022 – Dec 2022

St Louis, USA (Remote)

Founding Data Engineer

- Built the entire cloud data platform from scratch on AWS (S3, Lambda, Redshift, EventBridge) as the first data hire, establishing all data infrastructure, pipelines, and analytics for the startup.
- Designed ETL/ELT pipelines integrating 3 SaaS applications and APIs into a centralized data lake and warehouse, enabling a single source of truth for the business.
- Created Reverse ETL pipelines to sync key business metrics back to HubSpot CRM and Slack, automating manual reporting and saving 100s of hours.
- Developed executive BI dashboards tracking core SaaS metrics (MRR, ARR, LTV, CAC, Cohort Analysis), directly supporting investor/board reporting and data-driven decision-making.

<b>FGF Brands</b>	<b>Jun 2021 – Feb 2022</b>
<i>Senior BI Developer</i>	<i>Toronto, Canada</i>
<ul style="list-style-type: none"> <li>Managed and optimized company data marts, data models, and the enterprise data warehouse on SQL Server.</li> <li>Created operational dashboards for Supply Chain and Finance departments, improving visibility into key business processes.</li> <li>Engineered ETL/ELT pipelines to process IoT data from sensors, cameras, and manufacturing systems, automating production line monitoring and reducing manual inspection by over 50%.</li> <li>Architected a modern cloud-based data lake and warehouse solution on Microsoft Azure (ADLS, ADF, Synapse), migrating from on-premise infrastructure.</li> </ul>	

<b>Tamkeen Technologies</b>	<b>Aug 2019 – May 2021</b>
<i>Senior BI Specialist</i>	<i>Riyadh, Saudi Arabia</i>
<ul style="list-style-type: none"> <li>Developed a data product to parse and transform data from log files generated by 7 security tools and developed analytical dashboards to visualize the insights, built with Python, Spark and Unix.</li> <li>Designed a data warehouse for the internal Odoo ERP system covering HR, Procurement, and Finance departments, enabling monitoring of business processes and transactions.</li> <li>Built and automated dashboards and operational reports for the company executives that saved 100s of hours.</li> <li>Provided technical oversight and strategic direction for multiple national data warehouses supporting government products serving millions of users.</li> </ul>	

<b>SABIC (Saudi Basic Industries Corporation)</b>	<b>Oct 2011 – Jul 2019</b>
<i>Senior Data Consultant</i>	<i>Riyadh, Saudi Arabia</i>
<ul style="list-style-type: none"> <li>Implemented process mining using Celonis for order-to-cash and procure-to-pay workflows, and delivered 20+ analytical dashboards that improved supply chain team performance and shipment delivery accuracy by 25%.</li> <li>Drove a 4x increase (20% to 80%) in B2B E-Commerce platform adoption within one year by leading usage data analysis and delivering data-driven feature enhancements.</li> <li>Led a data migration from legacy systems into a centralized SAP ERP platform, building data quality monitoring dashboards to ensure data integrity throughout the transformation.</li> <li>Established SABIC's first enterprise-wide data governance and data quality program (DAMA-aligned) across Finance, HR, and Manufacturing, defining data profiling, quality scorecards, business glossary, and lineage analysis.</li> </ul>	

## Education

<b>King Abdullah University of Science and Technology (KAUST)</b>	<b>Aug 2010 – Sep 2011</b>
<i>M.Sc in Computer Science (Data Mining &amp; AI)</i>	<i>Thuwal, Saudi Arabia</i>
<b>King Fahd University of Petroleum and Minerals (KFUPM)</b>	<b>Feb 2006 – Aug 2010</b>
<i>B.Sc in Computer Engineering (Honors)</i>	<i>Dhahran, Saudi Arabia</i>

<b>Texas A&amp;M University</b>	<b>Aug 2008 – Dec 2008</b>
<i>Student Exchange Program, Computer Engineering</i>	<i>Texas, USA</i>

## Certificates

<b>Databricks:</b> Data Engineer (Associate, Professional), Data Analyst, GenAI Engineer	2025 – 2026
<b>Microsoft:</b> Fabric Data Engineer Associate	2026
<b>Google Cloud:</b> Professional Data Engineer	2024
<b>DAMA:</b> Certified Data Management Professional (CDMP)	2019
<b>TDWI:</b> Certified Business Intelligence Professional (CBIP)	2020
<b>The Open Group:</b> TOGAF Enterprise Architecture	2020
<b>PMI:</b> Project Management Professional (PMP)	2020

## Technical Skills

<b>Cloud Platforms:</b> Azure, AWS, GCP, Databricks, Microsoft Fabric	<b>Programming &amp; DevOps:</b> Python, SQL, C#, Git, Docker, Azure DevOps, CI/CD
<b>Lakehouse &amp; Warehousing:</b> Delta Lake, Databricks, Unity Catalog, Snowflake, Redshift, Synapse, BigQuery, SAP HANA	<b>BI &amp; Analytics:</b> Power BI, Tableau, Sigma Computing, Celonis Process Mining
<b>Data Engineering:</b> PySpark, Pandas, ADF, SSIS, Airflow, dbt	<b>Data Governance:</b> Atlan, SAP Information Steward, Unity Catalog, SAP Data Services
<b>GenAI &amp; ML:</b> Claude Code, Google ADK, LangChain, DSPy, Databricks Mosaic AI, MLflow	<b>Databases:</b> PostgreSQL, SQL Server, DynamoDB, Cloud SQL